

02_pandas_time_indexes

August 13, 2021

1 more on Pandas index

1.1 Credit:

this comes from Abernathys open book, which we will be looking at a lot! https://earth-env-data-science.github.io/lectures/core_python/python_fundamentals.html

1.2 Time Indexes

Indexes are very powerful. They label the data inside a pandas series or dataframe and let you intuitively work with the data. They are a big part of why Pandas is so useful. There are different indices for different types of data. Time Indexes are especially great!

```
[1]: # import pandas, etc
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```


2 dates

Python has a special datatype specific to dates. Dates get treated differently than nummbers or strings, and this feature lets you do lots of powerful timeseries analysis.

Below we make a special time series using `pd.date_range()` where we can speficy the start, end and frequency of points we want

```
[2]: two_years = pd.date_range(start='2014-01-01', end='2016-01-01', freq='D')
timeseries = pd.Series(np.sin(2 *np.pi *two_years.dayofyear / 365),
                        index=two_years)
timeseries.plot()
```

```
[2]: <AxesSubplot:>
```



What happened there? We made a range of times (data of a particular type, the pandas `datetime` type) using pandas `pd.date_range()`, then we used numpy (`np.sin()`) to make some fake data based on that time range.

I added the fake data, and the time range into a pandas `Series` called `timeseries`, which is like one column of a `DataFrame` using the function that creates series: `pd.Series()`, telling the function to use `two_years` as our index for our variable `timeseries`.


Because we used the special way pandas can create time indices, when we plotted the `Series` using `timeseries.plot()`, matplotlib knows we are talking about time and labels everything nicely.

3 indexing and slicing

Let's say we want to just get some time subset of our data. Pandas has an easy way to let us do that using the `.loc` notation we've seen before:

```
[3]: timeseries.loc['2015-01-01':'2015-07-01'].plot()
```

```
[3]: <AxesSubplot:>
```



The `TimeIndex` object has lots of useful attributes

```
[4]: timeseries.index.month
```

```
[4]: Int64Index([ 1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
...
               12, 12, 12, 12, 12, 12, 12, 12, 12,  1],
               dtype='int64', length=731)
```

```
[5]: timeseries.index.day
```

```
[5]: Int64Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10,
...
               23, 24, 25, 26, 27, 28, 29, 30, 31,  1],
               dtype='int64', length=731)
```

3.1 Reading real Data Files: NOAA Weather Station Data

In this example, we will use NOAA weather station data from <https://www.ncdc.noaa.gov/data-access/land-based-station-data>.

The details of files we are going to read are described in this [README file](#).

We have a text file on our hard drive called `data.txt`. **Examine it by opening in jupyterlab.** What do you see?

To read it into pandas, we will use the `read_csv` function. This function is incredibly complex

and powerful. You can use it to extract data from almost any text file. However, you need to understand how to use its various options. It's awesome, it will let you read data out of almost any text file. However to get the data in a useful format we are going to have to do a few things to it. In the next few steps we will see how to clean and wrangle the data into a good format for use.

3.1.1 Cleaning and wrangling messy data into a format you can use is an incredibly important skill to master

If we just read the data in with no options, we get something that is a bit of a mess:

With no options, this is what we get.

```
[6]: df = pd.read_csv('data.txt')  
  
print(df.shape)  
df.head()
```

(365, 1)

```
[6]: WBANNO LST_DATE CRX_VN LONGITUDE LATITUDE T_DAILY_MAX T_DAILY_MIN T_DAILY_MEAN  
T_DAILY_AVG P_DAILY_CALC SOLARAD_DAILY SUR_TEMP_DAILY_TYPE SUR_TEMP_DAILY_MAX  
SUR_TEMP_DAILY_MIN SUR_TEMP_DAILY_AVG RH_DAILY_MAX RH_DAILY_MIN RH_DAILY_AVG  
SOIL_MOISTURE_5_DAILY SOIL_MOISTURE_10_DAILY SOIL_MOISTURE_20_DAILY  
SOIL_MOISTURE_50_DAILY SOIL_MOISTURE_100_DAILY SOIL_TEMP_5_DAILY  
SOIL_TEMP_10_DAILY SOIL_TEMP_20_DAILY SOIL_TEMP_50_DAILY SOIL_TEMP_100_DAILY  
0 64756 20170101 2.422 -73.74 41.79 6.6 ...  
1 64756 20170102 2.422 -73.74 41.79 4.0 ...  
2 64756 20170103 2.422 -73.74 41.79 4.9 ...  
3 64756 20170104 2.422 -73.74 41.79 8.7 ...  
4 64756 20170105 2.422 -73.74 41.79 -0.5 ...
```

Pandas wasn't able to automatically sort out all the columns because, as the name suggests, `pd.read_csv()` expects the data to be in comma separated value (csv) format. Our data is just separated by spaces.

Fortunately we can put options into `pd.read_csv()` to tell it what the separation between data is by using the `sep=` keyword. This lets the function read the data correctly. The representation of space is `'\s+'`.

```
[7]: df = pd.read_csv('data.txt', sep='\s+')  
  
print(df.shape)  
df.head()
```

(365, 28)

```
[7]: WBANNO LST_DATE CRX_VN LONGITUDE LATITUDE T_DAILY_MAX T_DAILY_MIN \  
0 64756 20170101 2.422 -73.74 41.79 6.6 -5.4  
1 64756 20170102 2.422 -73.74 41.79 4.0 -6.8  
2 64756 20170103 2.422 -73.74 41.79 4.9 0.7  
3 64756 20170104 2.422 -73.74 41.79 8.7 -1.6  
4 64756 20170105 2.422 -73.74 41.79 -0.5 -4.6
```

	T_DAILY_MEAN	T_DAILY_AVG	P_DAILY_CALC	...	SOIL_MOISTURE_5_DAILY	\
0	0.6	2.2	0.0	...	-99.0	
1	-1.4	-1.2	0.0	...	-99.0	
2	2.8	2.7	13.1	...	-99.0	
3	3.6	3.5	1.3	...	-99.0	
4	-2.5	-2.8	0.0	...	-99.0	

	SOIL_MOISTURE_10_DAILY	SOIL_MOISTURE_20_DAILY	SOIL_MOISTURE_50_DAILY	\
0	-99.0	0.207	0.152	
1	-99.0	0.205	0.151	
2	-99.0	0.205	0.150	
3	-99.0	0.215	0.153	
4	-99.0	0.215	0.154	

	SOIL_MOISTURE_100_DAILY	SOIL_TEMP_5_DAILY	SOIL_TEMP_10_DAILY	\
0	0.175	-0.1	0.0	
1	0.173	-0.2	0.0	
2	0.173	-0.1	0.0	
3	0.174	-0.1	0.0	
4	0.177	-0.1	0.0	

	SOIL_TEMP_20_DAILY	SOIL_TEMP_50_DAILY	SOIL_TEMP_100_DAILY
0	0.6	1.5	3.4
1	0.6	1.5	3.3
2	0.5	1.5	3.3
3	0.5	1.5	3.2
4	0.5	1.4	3.1

[5 rows x 28 columns]

excellent, much better. now the columns are all separated out well.

if we looked at all the data we will see there are lots of -99 and -9999 values in the file. If we look closely, we will see there are lots of -99 and -9999 values in the file. The [README file](#) tells us that these are values used to represent missing data. Let's tell this to pandas. We do this using an argument specification `na_values = [list of bad data values]`. the `na` part stands for NaN (not a number) which will be filled in:

```
[8]: df = pd.read_csv('data.txt', sep='\s+', na_values=[-9999.0, -99.0])
df.head()
```

```
[8]:  WBANNO  LST_DATE  CRX_VN  LONGITUDE  LATITUDE  T_DAILY_MAX  T_DAILY_MIN  \
0    64756  20170101   2.422    -73.74    41.79         6.6        -5.4
1    64756  20170102   2.422    -73.74    41.79         4.0        -6.8
2    64756  20170103   2.422    -73.74    41.79         4.9         0.7
3    64756  20170104   2.422    -73.74    41.79         8.7        -1.6
4    64756  20170105   2.422    -73.74    41.79        -0.5        -4.6
```

	T_DAILY_MEAN	T_DAILY_AVG	P_DAILY_CALC	...	SOIL_MOISTURE_5_DAILY	\
0	0.6	2.2	0.0	...		NaN
1	-1.4	-1.2	0.0	...		NaN
2	2.8	2.7	13.1	...		NaN
3	3.6	3.5	1.3	...		NaN
4	-2.5	-2.8	0.0	...		NaN

	SOIL_MOISTURE_10_DAILY	SOIL_MOISTURE_20_DAILY	SOIL_MOISTURE_50_DAILY	\
0	NaN	0.207	0.152	
1	NaN	0.205	0.151	
2	NaN	0.205	0.150	
3	NaN	0.215	0.153	
4	NaN	0.215	0.154	

	SOIL_MOISTURE_100_DAILY	SOIL_TEMP_5_DAILY	SOIL_TEMP_10_DAILY	\
0	0.175	-0.1	0.0	
1	0.173	-0.2	0.0	
2	0.173	-0.1	0.0	
3	0.174	-0.1	0.0	
4	0.177	-0.1	0.0	

	SOIL_TEMP_20_DAILY	SOIL_TEMP_50_DAILY	SOIL_TEMP_100_DAILY
0	0.6	1.5	3.4
1	0.6	1.5	3.3
2	0.5	1.5	3.3
3	0.5	1.5	3.2
4	0.5	1.4	3.1

[5 rows x 28 columns]

ok, another good step. Now all the bad data is represented by NaN, which is something that pandas and numpy are good at dealing with.

you can see we are slowly building up a good dataframe. We are cleaning out all the warts in the data. This is something you will do over and over!

Let's check out what is in our Dataframe:

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 28 columns):
#   Column              Non-Null Count  Dtype
---  -
0   WBANNO              365 non-null    int64
1   LST_DATE            365 non-null    int64
2   CRX_VN              365 non-null    float64
```

```

3  LONGITUDE          365 non-null    float64
4  LATITUDE           365 non-null    float64
5  T_DAILY_MAX        364 non-null    float64
6  T_DAILY_MIN        364 non-null    float64
7  T_DAILY_MEAN       364 non-null    float64
8  T_DAILY_AVG        364 non-null    float64
9  P_DAILY_CALC       364 non-null    float64
10 SOLARAD_DAILY      364 non-null    float64
11 SUR_TEMP_DAILY_TYPE 365 non-null    object
12 SUR_TEMP_DAILY_MAX  364 non-null    float64
13 SUR_TEMP_DAILY_MIN  364 non-null    float64
14 SUR_TEMP_DAILY_AVG  364 non-null    float64
15 RH_DAILY_MAX        364 non-null    float64
16 RH_DAILY_MIN        364 non-null    float64
17 RH_DAILY_AVG        364 non-null    float64
18 SOIL_MOISTURE_5_DAILY 317 non-null    float64
19 SOIL_MOISTURE_10_DAILY 317 non-null    float64
20 SOIL_MOISTURE_20_DAILY 336 non-null    float64
21 SOIL_MOISTURE_50_DAILY 364 non-null    float64
22 SOIL_MOISTURE_100_DAILY 359 non-null    float64
23 SOIL_TEMP_5_DAILY   364 non-null    float64
24 SOIL_TEMP_10_DAILY  364 non-null    float64
25 SOIL_TEMP_20_DAILY  364 non-null    float64
26 SOIL_TEMP_50_DAILY  364 non-null    float64
27 SOIL_TEMP_100_DAILY 364 non-null    float64
dtypes: float64(25), int64(2), object(1)
memory usage: 80.0+ KB

```

One problem here is that pandas did not recognize the LDT_DATE column as a date. Let's help it.

```

[10]: df = pd.read_csv('data.txt', sep='\s+',
                      na_values=[-9999.0, -99.0],
                      parse_dates=[1])
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 28 columns):
#   Column              Non-Null Count  Dtype
---  -
0   WBANNO              365 non-null    int64
1   LST_DATE            365 non-null    datetime64[ns]
2   CRX_VN              365 non-null    float64
3   LONGITUDE           365 non-null    float64
4   LATITUDE            365 non-null    float64
5   T_DAILY_MAX         364 non-null    float64
6   T_DAILY_MIN         364 non-null    float64
7   T_DAILY_MEAN        364 non-null    float64
8   T_DAILY_AVG         364 non-null    float64

```

```

9   P_DAILY_CALC          364 non-null    float64
10  SOLARAD_DAILY          364 non-null    float64
11  SUR_TEMP_DAILY_TYPE    365 non-null    object
12  SUR_TEMP_DAILY_MAX     364 non-null    float64
13  SUR_TEMP_DAILY_MIN     364 non-null    float64
14  SUR_TEMP_DAILY_AVG     364 non-null    float64
15  RH_DAILY_MAX           364 non-null    float64
16  RH_DAILY_MIN           364 non-null    float64
17  RH_DAILY_AVG           364 non-null    float64
18  SOIL_MOISTURE_5_DAILY   317 non-null    float64
19  SOIL_MOISTURE_10_DAILY  317 non-null    float64
20  SOIL_MOISTURE_20_DAILY  336 non-null    float64
21  SOIL_MOISTURE_50_DAILY  364 non-null    float64
22  SOIL_MOISTURE_100_DAILY 359 non-null    float64
23  SOIL_TEMP_5_DAILY      364 non-null    float64
24  SOIL_TEMP_10_DAILY     364 non-null    float64
25  SOIL_TEMP_20_DAILY     364 non-null    float64
26  SOIL_TEMP_50_DAILY     364 non-null    float64
27  SOIL_TEMP_100_DAILY    364 non-null    float64
dtypes: datetime64[ns](1), float64(25), int64(1), object(1)
memory usage: 80.0+ KB

```

now we see that the LST_DATE column is the special datetime64 type of data that lets pandas do all the cool stuff with time.

We are one step closer to a good data set!

the last step: we want to use the date as the index for our dataframe. It's the timestamp that ties all this data together. We can tell pandas to do this too by setting the index of our dataframe to that column

```
[11]: df = df.set_index('LST_DATE')
      df.head()
```

```
[11]:
```

	WBANNO	CRX_VN	LONGITUDE	LATITUDE	T_DAILY_MAX	T_DAILY_MIN	\
LST_DATE							
2017-01-01	64756	2.422	-73.74	41.79	6.6	-5.4	
2017-01-02	64756	2.422	-73.74	41.79	4.0	-6.8	
2017-01-03	64756	2.422	-73.74	41.79	4.9	0.7	
2017-01-04	64756	2.422	-73.74	41.79	8.7	-1.6	
2017-01-05	64756	2.422	-73.74	41.79	-0.5	-4.6	

	T_DAILY_MEAN	T_DAILY_AVG	P_DAILY_CALC	SOLARAD_DAILY	...	\
LST_DATE						
2017-01-01	0.6	2.2	0.0	8.68	...	
2017-01-02	-1.4	-1.2	0.0	2.08	...	
2017-01-03	2.8	2.7	13.1	0.68	...	
2017-01-04	3.6	3.5	1.3	2.85	...	
2017-01-05	-2.5	-2.8	0.0	4.90	...	

	SOIL_MOISTURE_5_DAILY	SOIL_MOISTURE_10_DAILY \
LST_DATE		
2017-01-01	NaN	NaN
2017-01-02	NaN	NaN
2017-01-03	NaN	NaN
2017-01-04	NaN	NaN
2017-01-05	NaN	NaN

	SOIL_MOISTURE_20_DAILY	SOIL_MOISTURE_50_DAILY \
LST_DATE		
2017-01-01	0.207	0.152
2017-01-02	0.205	0.151
2017-01-03	0.205	0.150
2017-01-04	0.215	0.153
2017-01-05	0.215	0.154

	SOIL_MOISTURE_100_DAILY	SOIL_TEMP_5_DAILY	SOIL_TEMP_10_DAILY \
LST_DATE			
2017-01-01	0.175	-0.1	0.0
2017-01-02	0.173	-0.2	0.0
2017-01-03	0.173	-0.1	0.0
2017-01-04	0.174	-0.1	0.0
2017-01-05	0.177	-0.1	0.0

	SOIL_TEMP_20_DAILY	SOIL_TEMP_50_DAILY	SOIL_TEMP_100_DAILY
LST_DATE			
2017-01-01	0.6	1.5	3.4
2017-01-02	0.6	1.5	3.3
2017-01-03	0.5	1.5	3.3
2017-01-04	0.5	1.5	3.2
2017-01-05	0.5	1.4	3.1

[5 rows x 27 columns]

Now we can take advantage of all the cool time stuff that pandas can do.

Like let's use indexing to look at all the data from one day

```
[12]: df.loc['2017-08-07']
```

```
[12]: WBANNO          64756
      CRX_VN          2.422
      LONGITUDE      -73.74
      LATITUDE       41.79
      T_DAILY_MAX     19.3
      T_DAILY_MIN     12.3
      T_DAILY_MEAN    15.8
```

```

T_DAILY_AVG          16.3
P_DAILY_CALC          4.9
SOLARAD_DAILY        3.93
SUR_TEMP_DAILY_TYPE   C
SUR_TEMP_DAILY_MAX    22.3
SUR_TEMP_DAILY_MIN    11.9
SUR_TEMP_DAILY_AVG    17.7
RH_DAILY_MAX          94.7
RH_DAILY_MIN          76.4
RH_DAILY_AVG          89.5
SOIL_MOISTURE_5_DAILY 0.148
SOIL_MOISTURE_10_DAILY 0.113
SOIL_MOISTURE_20_DAILY 0.094
SOIL_MOISTURE_50_DAILY 0.114
SOIL_MOISTURE_100_DAILY 0.151
SOIL_TEMP_5_DAILY     21.4
SOIL_TEMP_10_DAILY    21.7
SOIL_TEMP_20_DAILY    22.1
SOIL_TEMP_50_DAILY    22.2
SOIL_TEMP_100_DAILY   21.5
Name: 2017-08-07 00:00:00, dtype: object

```

Or use slicing to get a range:

```
[13]: df.loc['2017-07-01':'2017-07-31']
```

```

[13]:      WBANNO  CRX_VN  LONGITUDE  LATITUDE  T_DAILY_MAX  T_DAILY_MIN  \
LST_DATE
2017-07-01  64756  2.422    -73.74    41.79         28.0         19.7
2017-07-02  64756  2.422    -73.74    41.79         29.8         18.4
2017-07-03  64756  2.422    -73.74    41.79         28.3         15.0
2017-07-04  64756  2.422    -73.74    41.79         26.8         12.6
2017-07-05  64756  2.422    -73.74    41.79         28.0         11.9
2017-07-06  64756  2.422    -73.74    41.79         25.7         14.3
2017-07-07  64756  2.422    -73.74    41.79         25.8         16.8
2017-07-08  64756  2.422    -73.74    41.79         29.0         15.3
2017-07-09  64756  2.422    -73.74    41.79         26.3         10.9
2017-07-10  64756  2.422    -73.74    41.79         27.6         11.8
2017-07-11  64756  2.422    -73.74    41.79         27.4         19.2
2017-07-12  64756  2.422    -73.74    41.79         29.4         18.5
2017-07-13  64756  2.422    -73.74    41.79         29.5         18.3
2017-07-14  64756  2.422    -73.74    41.79         18.5         15.9
2017-07-15  64756  2.422    -73.74    41.79         26.6         16.5
2017-07-16  64756  2.422    -73.74    41.79         27.9         13.3
2017-07-17  64756  2.422    -73.74    41.79         29.2         16.1
2017-07-18  64756  2.422    -73.74    41.79         30.3         19.3
2017-07-19  64756  2.422    -73.74    41.79         31.2         19.1
2017-07-20  64756  2.422    -73.74    41.79         31.8         16.6

```

2017-07-21	64756	2.422	-73.74	41.79	30.6	16.6
2017-07-22	64756	2.422	-73.74	41.79	27.7	15.6
2017-07-23	64756	2.422	-73.74	41.79	26.4	18.5
2017-07-24	64756	2.422	-73.74	41.79	19.4	14.8
2017-07-25	64756	2.422	-73.74	41.79	18.6	13.7
2017-07-26	64756	2.422	-73.74	41.79	24.7	11.2
2017-07-27	64756	2.422	-73.74	41.79	24.2	15.2
2017-07-28	64756	2.422	-73.74	41.79	26.5	16.9
2017-07-29	64756	2.422	-73.74	41.79	24.2	10.4
2017-07-30	64756	2.422	-73.74	41.79	25.5	8.2
2017-07-31	64756	2.422	-73.74	41.79	29.4	10.1

	T_DAILY_MEAN	T_DAILY_AVG	P_DAILY_CALC	SOLARAD_DAILY	...	\
LST_DATE					...	
2017-07-01	23.9	23.8	0.2	19.28	...	
2017-07-02	24.1	23.7	4.0	27.67	...	
2017-07-03	21.7	21.4	0.0	27.08	...	
2017-07-04	19.7	20.0	0.0	29.45	...	
2017-07-05	20.0	20.7	0.0	26.90	...	
2017-07-06	20.0	20.3	0.0	19.03	...	
2017-07-07	21.3	20.0	11.5	13.88	...	
2017-07-08	22.1	21.5	0.0	21.92	...	
2017-07-09	18.6	19.4	0.0	29.72	...	
2017-07-10	19.7	21.3	0.0	23.67	...	
2017-07-11	23.3	22.6	8.5	17.79	...	
2017-07-12	23.9	23.1	1.9	16.27	...	
2017-07-13	23.9	23.4	23.3	13.61	...	
2017-07-14	17.2	17.5	4.1	5.36	...	
2017-07-15	21.5	21.0	0.8	21.13	...	
2017-07-16	20.6	21.0	0.0	27.03	...	
2017-07-17	22.6	22.9	0.0	20.47	...	
2017-07-18	24.8	24.7	0.0	24.99	...	
2017-07-19	25.1	25.0	0.0	27.69	...	
2017-07-20	24.2	23.4	0.7	21.53	...	
2017-07-21	23.6	23.6	0.0	25.55	...	
2017-07-22	21.7	21.2	0.5	16.04	...	
2017-07-23	22.5	22.2	0.0	19.03	...	
2017-07-24	17.1	16.7	29.2	9.10	...	
2017-07-25	16.2	16.2	0.0	7.35	...	
2017-07-26	18.0	18.3	0.0	22.22	...	
2017-07-27	19.7	19.5	0.0	8.28	...	
2017-07-28	21.7	20.9	0.0	21.06	...	
2017-07-29	17.3	18.1	0.0	21.28	...	
2017-07-30	16.8	17.3	0.0	27.68	...	
2017-07-31	19.7	20.1	0.0	25.49	...	

SOIL_MOISTURE_5_DAILY SOIL_MOISTURE_10_DAILY \

LST_DATE		
2017-07-01	0.157	0.136
2017-07-02	0.146	0.135
2017-07-03	0.141	0.132
2017-07-04	0.131	0.126
2017-07-05	0.116	0.114
2017-07-06	0.105	0.104
2017-07-07	0.114	0.100
2017-07-08	0.130	0.106
2017-07-09	0.119	0.103
2017-07-10	0.105	0.096
2017-07-11	0.106	0.093
2017-07-12	0.108	0.094
2017-07-13	0.134	0.110
2017-07-14	0.194	0.151
2017-07-15	0.190	0.163
2017-07-16	0.171	0.154
2017-07-17	0.155	0.143
2017-07-18	0.142	0.132
2017-07-19	0.126	0.118
2017-07-20	0.111	0.103
2017-07-21	0.100	0.093
2017-07-22	0.092	0.086
2017-07-23	0.087	0.082
2017-07-24	0.145	0.118
2017-07-25	0.167	0.133
2017-07-26	0.155	0.128
2017-07-27	0.144	0.122
2017-07-28	0.137	0.117
2017-07-29	0.126	0.108
2017-07-30	0.113	0.099
2017-07-31	0.101	0.090

LST_DATE	SOIL_MOISTURE_20_DAILY	SOIL_MOISTURE_50_DAILY \
2017-07-01	0.144	0.129
2017-07-02	0.143	0.129
2017-07-03	0.139	0.128
2017-07-04	0.136	0.126
2017-07-05	0.131	0.125
2017-07-06	0.126	0.124
2017-07-07	0.123	0.123
2017-07-08	0.122	0.123
2017-07-09	0.119	0.121
2017-07-10	0.113	0.120
2017-07-11	0.110	0.120
2017-07-12	0.108	0.118

2017-07-13	0.108	0.118
2017-07-14	0.114	0.120
2017-07-15	0.119	0.122
2017-07-16	0.123	0.123
2017-07-17	0.124	0.122
2017-07-18	0.122	0.122
2017-07-19	0.118	0.122
2017-07-20	0.114	0.121
2017-07-21	0.108	0.120
2017-07-22	0.104	0.119
2017-07-23	0.100	0.118
2017-07-24	0.102	0.117
2017-07-25	0.107	0.116
2017-07-26	0.108	0.118
2017-07-27	0.109	0.118
2017-07-28	0.110	0.119
2017-07-29	0.108	0.118
2017-07-30	0.104	0.117
2017-07-31	0.099	0.116

	SOIL_MOISTURE_100_DAILY	SOIL_TEMP_5_DAILY	SOIL_TEMP_10_DAILY \
LST_DATE			
2017-07-01	0.163	25.7	25.4
2017-07-02	0.162	26.8	26.4
2017-07-03	0.162	26.4	26.3
2017-07-04	0.161	25.9	25.8
2017-07-05	0.161	25.3	25.3
2017-07-06	0.160	24.7	24.7
2017-07-07	0.160	24.2	24.2
2017-07-08	0.159	25.5	25.3
2017-07-09	0.158	24.8	24.8
2017-07-10	0.158	24.7	24.7
2017-07-11	0.157	25.6	25.4
2017-07-12	0.157	25.8	25.6
2017-07-13	0.156	25.7	25.7
2017-07-14	0.155	23.0	23.3
2017-07-15	0.155	24.6	24.4
2017-07-16	0.155	25.4	25.3
2017-07-17	0.156	25.7	25.6
2017-07-18	0.156	27.0	26.7
2017-07-19	0.156	27.6	27.4
2017-07-20	0.156	27.0	27.0
2017-07-21	0.155	27.1	27.0
2017-07-22	0.156	25.9	26.1
2017-07-23	0.155	26.0	26.0
2017-07-24	0.154	23.1	23.6
2017-07-25	0.153	21.9	22.2

2017-07-26	0.152	22.9	23.0
2017-07-27	0.154	22.5	22.7
2017-07-28	0.154	24.1	24.1
2017-07-29	0.154	23.3	23.6
2017-07-30	0.154	22.8	23.0
2017-07-31	0.153	23.8	23.8

LST_DATE	SOIL_TEMP_20_DAILY	SOIL_TEMP_50_DAILY	SOIL_TEMP_100_DAILY
2017-07-01	23.7	21.9	19.9
2017-07-02	24.5	22.3	20.1
2017-07-03	24.8	22.8	20.3
2017-07-04	24.6	22.9	20.6
2017-07-05	24.2	22.8	20.7
2017-07-06	23.9	22.7	20.9
2017-07-07	23.4	22.4	20.8
2017-07-08	23.9	22.4	20.8
2017-07-09	23.8	22.5	20.8
2017-07-10	23.6	22.5	20.9
2017-07-11	24.1	22.6	20.9
2017-07-12	24.2	22.8	21.0
2017-07-13	24.4	23.0	21.0
2017-07-14	23.4	22.9	21.2
2017-07-15	23.2	22.2	21.2
2017-07-16	23.9	22.6	21.1
2017-07-17	24.4	22.9	21.2
2017-07-18	24.9	23.2	21.3
2017-07-19	25.6	23.7	21.5
2017-07-20	25.6	24.0	21.7
2017-07-21	25.5	24.0	21.9
2017-07-22	25.3	24.1	22.0
2017-07-23	24.9	23.8	22.1
2017-07-24	23.9	23.5	22.1
2017-07-25	22.4	22.5	21.9
2017-07-26	22.3	22.0	21.7
2017-07-27	22.4	22.0	21.4
2017-07-28	22.8	22.0	21.3
2017-07-29	23.0	22.2	21.3
2017-07-30	22.4	22.0	21.3
2017-07-31	22.7	21.9	21.2

[31 rows x 27 columns]

3.1.2 Quick Statistics

```
[14]: df.describe()
```

```
[14]:
```

	WBANNO	CRX_VN	LONGITUDE	LATITUDE	T_DAILY_MAX	\
count	365.0	365.000000	3.650000e+02	3.650000e+02	364.000000	
mean	64756.0	2.470767	-7.374000e+01	4.179000e+01	15.720055	
std	0.0	0.085997	5.265234e-13	3.842198e-13	10.502087	
min	64756.0	2.422000	-7.374000e+01	4.179000e+01	-12.300000	
25%	64756.0	2.422000	-7.374000e+01	4.179000e+01	6.900000	
50%	64756.0	2.422000	-7.374000e+01	4.179000e+01	17.450000	
75%	64756.0	2.422000	-7.374000e+01	4.179000e+01	24.850000	
max	64756.0	2.622000	-7.374000e+01	4.179000e+01	33.400000	

	T_DAILY_MIN	T_DAILY_MEAN	T_DAILY_AVG	P_DAILY_CALC	SOLARAD_DAILY	\
count	364.000000	364.000000	364.000000	364.000000	364.000000	
mean	4.037912	9.876374	9.990110	2.797802	13.068187	
std	9.460676	9.727451	9.619168	7.238628	7.953074	
min	-21.800000	-17.000000	-16.700000	0.000000	0.100000	
25%	-2.775000	2.100000	2.275000	0.000000	6.225000	
50%	4.350000	10.850000	11.050000	0.000000	12.865000	
75%	11.900000	18.150000	18.450000	1.400000	19.740000	
max	20.700000	25.700000	26.700000	65.700000	29.910000	

	...	SOIL_MOISTURE_5_DAILY	SOIL_MOISTURE_10_DAILY	\
count	...	317.000000	317.000000	
mean	...	0.183804	0.181000	
std	...	0.047493	0.052697	
min	...	0.075000	0.074000	
25%	...	0.148000	0.137000	
50%	...	0.192000	0.198000	
75%	...	0.221000	0.219000	
max	...	0.294000	0.321000	

	SOIL_MOISTURE_20_DAILY	SOIL_MOISTURE_50_DAILY	\
count	336.000000	364.000000	
mean	0.156533	0.138286	
std	0.042775	0.019207	
min	0.069000	0.100000	
25%	0.118000	0.118000	
50%	0.169000	0.147000	
75%	0.188000	0.152250	
max	0.231000	0.170000	

	SOIL_MOISTURE_100_DAILY	SOIL_TEMP_5_DAILY	SOIL_TEMP_10_DAILY	\
count	359.000000	364.000000	364.000000	
mean	0.162844	12.344231	12.308516	

std	0.013814	9.367742	9.350273
min	0.128000	-0.700000	-0.400000
25%	0.155000	2.275000	2.075000
50%	0.166000	13.300000	13.350000
75%	0.173000	21.025000	21.125000
max	0.192000	27.600000	27.400000

	SOIL_TEMP_20_DAILY	SOIL_TEMP_50_DAILY	SOIL_TEMP_100_DAILY
count	364.000000	364.000000	364.000000
mean	12.060989	11.960989	11.971978
std	8.760899	8.082595	7.170197
min	0.200000	0.900000	1.900000
25%	2.575000	3.300000	4.100000
50%	13.100000	12.850000	11.650000
75%	20.400000	19.800000	19.325000
max	25.600000	24.100000	22.100000

[8 rows x 26 columns]


3.1.3 Plotting Values

We can now quickly make plots of the data

Pandas is very “time aware”:

```
[15]: df.T_DAILY_MEAN.plot()
```

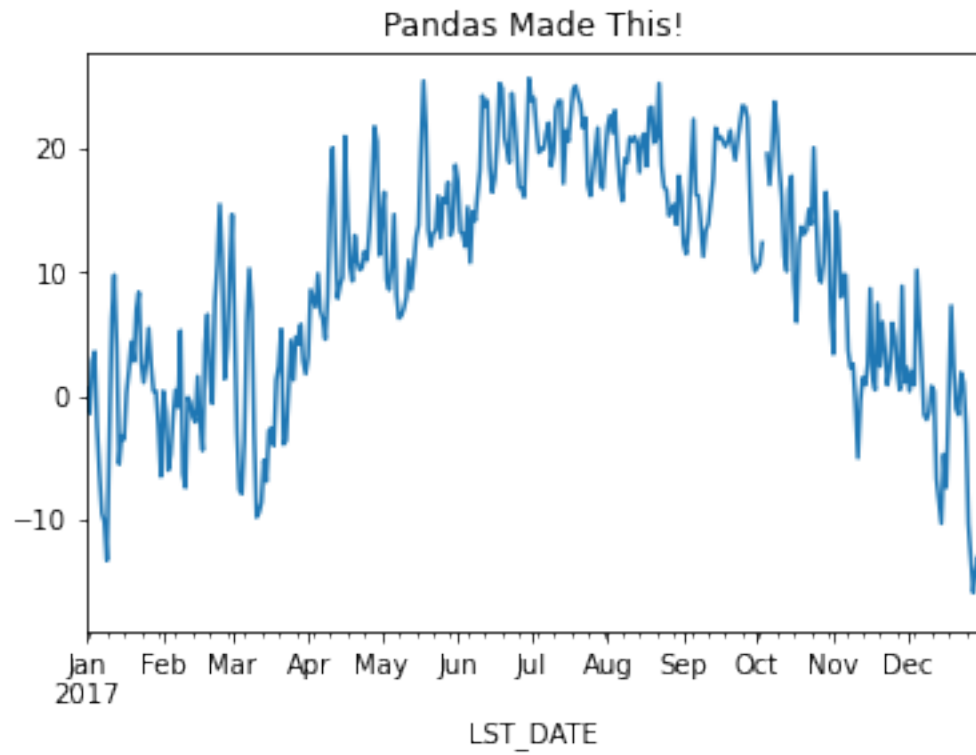
```
[15]: <AxesSubplot:xlabel='LST_DATE'>
```

Note: we could also manually create an axis and plot into it.

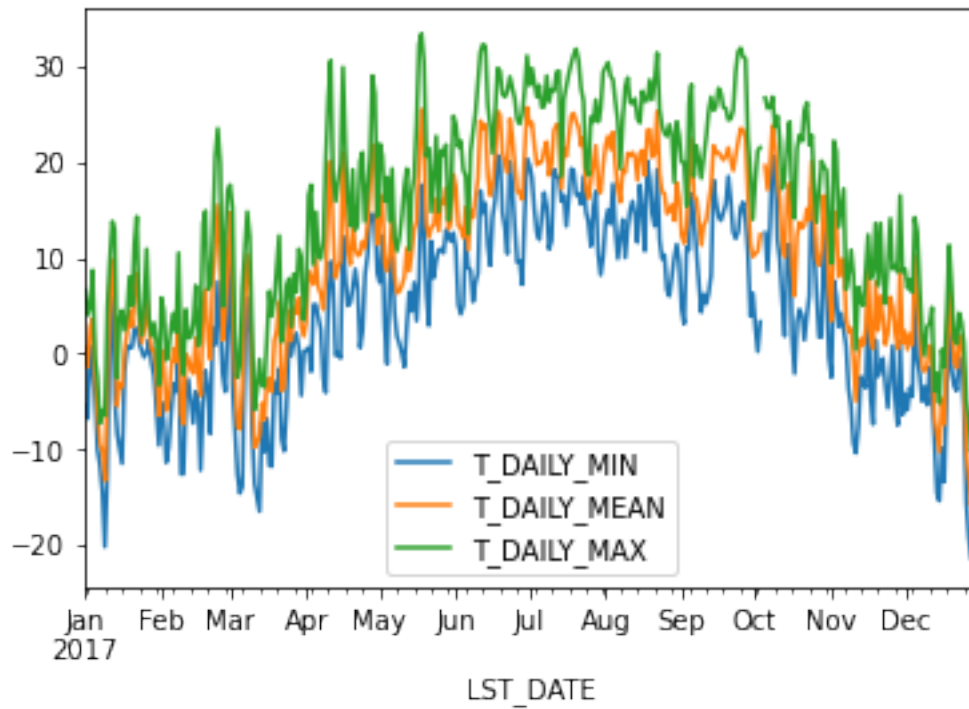
```
[16]: fig, ax = plt.subplots()
      df.T_DAILY_MEAN.plot(ax=ax)
      ax.set_title('Pandas Made This!')
```

```
[16]: Text(0.5, 1.0, 'Pandas Made This!')
```



```
[17]: df[['T_DAILY_MIN', 'T_DAILY_MEAN', 'T_DAILY_MAX']].plot()
```

```
[17]: <AxesSubplot:xlabel='LST_DATE'>
```

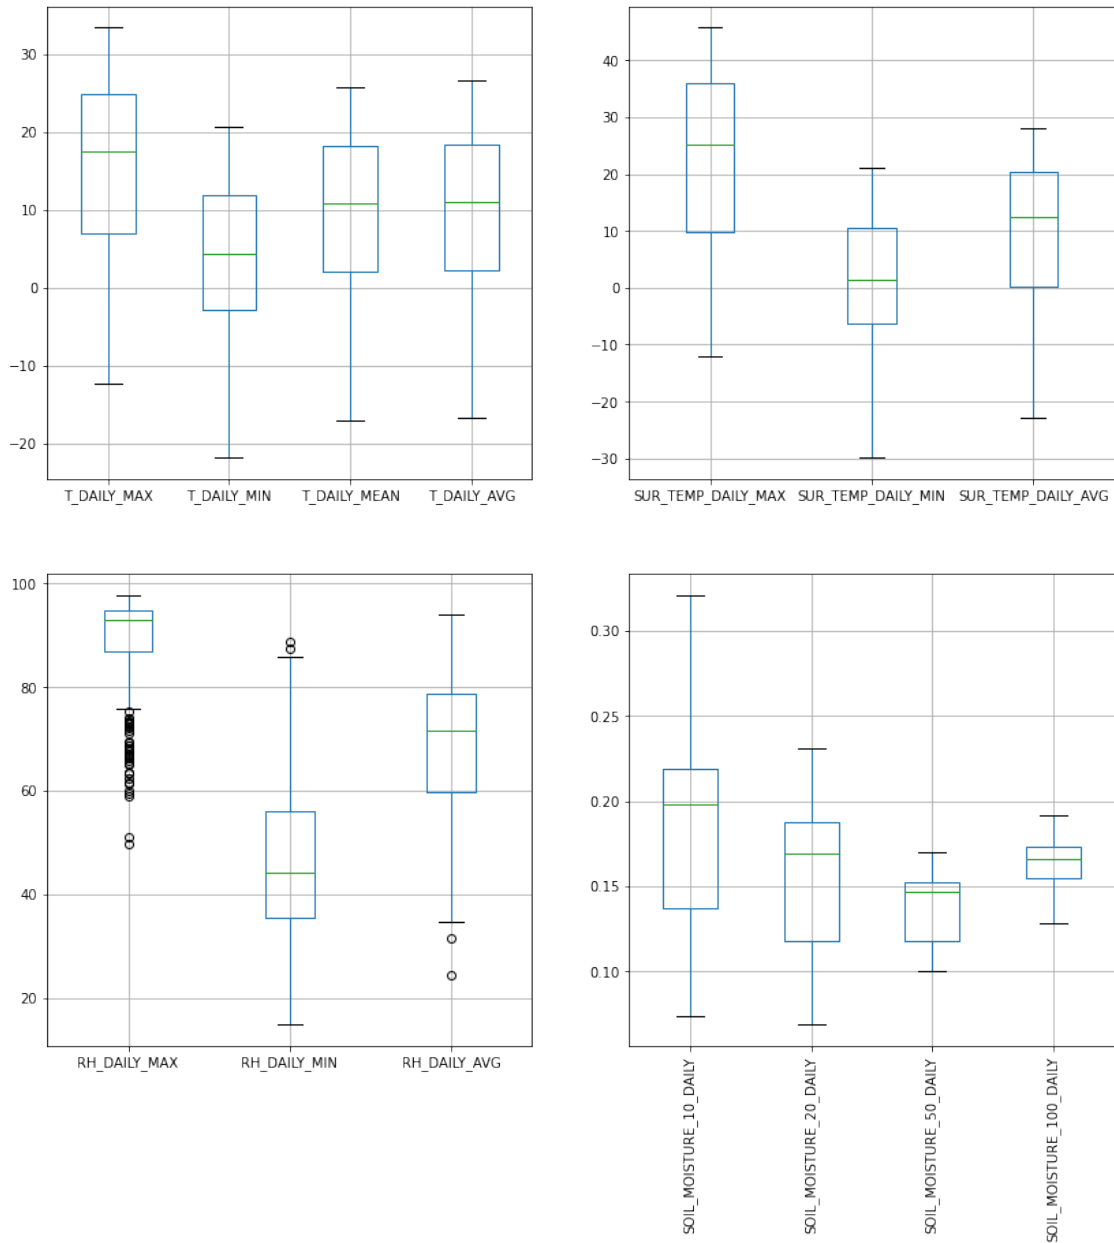


We can do some statistical plots too:

```
[18]: fig, ax = plt.subplots(ncols=2, nrows=2, figsize=(14,14))

df.iloc[:, 4:8].boxplot(ax=ax[0,0])
df.iloc[:, 10:14].boxplot(ax=ax[0,1])
df.iloc[:, 14:17].boxplot(ax=ax[1,0])
df.iloc[:, 18:22].boxplot(ax=ax[1,1])

ax[1, 1].set_xticklabels(ax[1, 1].get_xticklabels(), rotation=90);
```



4 Resampling

Pandas understands how to do all sorts of time related things. It has many methods that let us manipulate timeseries in useful ways.

For example there is a method called resampling that oranzizes the data based on a different time bin than the original data.

Let's resample from daily to monthly data:

```
[19]: # monthly reampler object
rs_obj = df.resample('MS')
rs_obj
```

```
[19]: <pandas.core.resample.DatetimeIndexResampler object at 0x7f9170d67160>
```

The output of the `.resample()` function is an object where the data has been stored into bins where all the data in a bin come from the same month. We then need to do some calculation on that bin to get a value. This is a `groupby` - type operation. we can use this to get monthly-means of all the data

```
[20]: rs_obj.mean()
```

```
[20]:
```

	WBANNO	CRX_VN	LONGITUDE	LATITUDE	T_DAILY_MAX	T_DAILY_MIN	\
LST_DATE							
2017-01-01	64756.0	2.422000	-73.74	41.79	3.945161	-3.993548	
2017-02-01	64756.0	2.422000	-73.74	41.79	7.246429	-4.360714	
2017-03-01	64756.0	2.422000	-73.74	41.79	5.164516	-5.335484	
2017-04-01	64756.0	2.422000	-73.74	41.79	17.813333	5.170000	
2017-05-01	64756.0	2.422000	-73.74	41.79	19.151613	7.338710	
2017-06-01	64756.0	2.422000	-73.74	41.79	25.423333	12.176667	
2017-07-01	64756.0	2.422000	-73.74	41.79	26.912903	15.183871	
2017-08-01	64756.0	2.422000	-73.74	41.79	25.741935	12.954839	
2017-09-01	64756.0	2.422000	-73.74	41.79	24.186667	11.300000	
2017-10-01	64756.0	2.602645	-73.74	41.79	21.043333	7.150000	
2017-11-01	64756.0	2.622000	-73.74	41.79	10.346667	-2.093333	
2017-12-01	64756.0	2.622000	-73.74	41.79	1.496774	-7.412903	

	T_DAILY_MEAN	T_DAILY_AVG	P_DAILY_CALC	SOLARAD_DAILY	...	\
LST_DATE						
2017-01-01	-0.025806	0.038710	3.090323	4.690000	...	
2017-02-01	1.442857	1.839286	2.414286	10.364286	...	
2017-03-01	-0.090323	0.167742	3.970968	13.113548	...	
2017-04-01	11.493333	11.540000	2.300000	14.645000	...	
2017-05-01	13.229032	13.638710	4.141935	16.519677	...	
2017-06-01	18.796667	18.986667	3.743333	21.655000	...	
2017-07-01	21.048387	20.993548	2.732258	20.566129	...	
2017-08-01	19.351613	19.477419	2.758065	18.360000	...	
2017-09-01	17.746667	17.463333	1.893333	15.154667	...	
2017-10-01	14.100000	13.976667	3.500000	10.395000	...	
2017-11-01	4.120000	4.336667	0.826667	6.723333	...	
2017-12-01	-2.967742	-2.838710	2.109677	4.474194	...	

	SOIL_MOISTURE_5_DAILY	SOIL_MOISTURE_10_DAILY	\
LST_DATE			
2017-01-01	0.236900	0.248300	
2017-02-01	0.226333	0.243000	
2017-03-01	0.218033	0.229267	

2017-04-01	0.199733	0.210300
2017-05-01	0.206613	0.210935
2017-06-01	0.185167	0.184300
2017-07-01	0.131226	0.115774
2017-08-01	0.143871	0.122258
2017-09-01	0.145167	0.139633
2017-10-01	0.140567	0.131467
2017-11-01	0.215433	0.211233
2017-12-01	0.231536	0.226143

	SOIL_MOISTURE_20_DAILY	SOIL_MOISTURE_50_DAILY \
LST_DATE		
2017-01-01	0.204550	0.152806
2017-02-01	0.207545	0.152857
2017-03-01	0.196258	0.153484
2017-04-01	0.190667	0.151000
2017-05-01	0.185613	0.147710
2017-06-01	0.173167	0.142533
2017-07-01	0.116613	0.121032
2017-08-01	0.105452	0.115290
2017-09-01	0.117267	0.112167
2017-10-01	0.084967	0.105667
2017-11-01	0.167067	0.149800
2017-12-01	0.177581	0.155516

	SOIL_MOISTURE_100_DAILY	SOIL_TEMP_5_DAILY	SOIL_TEMP_10_DAILY \
LST_DATE			
2017-01-01	0.175194	0.209677	0.267742
2017-02-01	0.175786	1.125000	1.100000
2017-03-01	0.174548	2.122581	2.161290
2017-04-01	0.172400	11.066667	10.666667
2017-05-01	0.170000	16.454839	16.290323
2017-06-01	0.167000	22.350000	22.166667
2017-07-01	0.156677	24.993548	24.980645
2017-08-01	0.151034	23.374194	23.519355
2017-09-01	0.141926	20.256667	20.386667
2017-10-01	0.133367	16.133333	16.186667
2017-11-01	0.164367	7.230000	7.190000
2017-12-01	0.169161	2.222581	2.187097

	SOIL_TEMP_20_DAILY	SOIL_TEMP_50_DAILY	SOIL_TEMP_100_DAILY
LST_DATE			
2017-01-01	0.696774	1.438710	2.877419
2017-02-01	1.192857	1.492857	2.367857
2017-03-01	2.345161	2.700000	3.387097
2017-04-01	9.636667	8.426667	6.903333
2017-05-01	15.361290	14.270968	12.696774

2017-06-01	20.880000	19.370000	17.333333
2017-07-01	23.925806	22.745161	21.164516
2017-08-01	22.848387	22.193548	21.377419
2017-09-01	19.966667	19.766667	19.530000
2017-10-01	16.320000	16.836667	17.470000
2017-11-01	8.060000	9.543333	11.746667
2017-12-01	2.916129	4.190323	6.303226

[12 rows x 26 columns]

we can chain all these commands together to plot the monthly mean of average, high and low temperatures:

```
[21]: df_mm = df.resample('MS').mean()
df_mm[['T_DAILY_MIN', 'T_DAILY_MEAN', 'T_DAILY_MAX']].plot()
```

```
[21]: <AxesSubplot:xlabel='LST_DATE'>
```

